


Question 1

```
import jax
from matplotlib import pyplot as plt
import time

def get_data_generator(batchsize):
    def getbatch(key):
        key, subkey = jax.random.split(key)
        context = 10*jax.random.uniform(subkey, [batchsize])
        key, subkey = jax.random.split(key)
        output = 3 + 10 * context - 0.7 * context**2 +
        jax.random.normal(subkey, [batchsize])
        return context, output
    return getbatch

gen = get_data_generator(100)

key = jax.random.PRNGKey(0)
context, output = gen(key)
plt.plot(context, output, 'o')
plt.xlabel('context')
plt.ylabel('output')
plt.show()
```



Question 2

```
import time
import jax
from jax import flatten_util
from jax import numpy as jnp
from typing import Callable
def signed_gradient_descent(net: Callable, loss: Callable, getbatch:
Callable, max_iter: int, learning_rates: list[int], *params):
    assert isinstance(net,Callable)
    assert isinstance(loss,Callable)
    assert isinstance(getbatch,Callable)
    assert isinstance(max_iter,int)

    key = jax.random.PRNGKey(0)
    w, unflatten = jax.flatten_util.ravel_pytree(params)

    # batch predictive network over context, but not params
    batched_net = jax.vmap(net,[0]+[None]*len(params))
    # batch loss over both context and predictions
    batched_loss = jax.vmap(loss)

    def l(key,w):
        params = unflatten(w)
        context, next = getbatch(key)
        pred = batched_net(context,*params)
        return jnp.mean(batched_loss(pred, next))

    fun = jax.value_and_grad(l,1)

    t0 = time.time()

    avg_grad = 0
    avg_loss = 0
    smooth_loss = 0

    print(" iter      l.r.      loss (smooth)      (avg)      time")
    print("-----")

    for i in range(max_iter):
        n = int(i * len(learning_rates) / max_iter)

        key, subkey = jax.random.split(key)
        loss, grad = fun(subkey, w)
        avg_loss = (i*avg_loss + loss)/(i+1)
        alpha = max(.01,1/(i+1))
        smooth_loss = alpha*loss + (1-alpha) * smooth_loss

        if i % (max_iter // 20) == 0:
            print(f"{i: >5} {learning_rates[n]:8.5f} {loss:8.5f}
```

```

{smooth_loss:8.5f} {avg_loss:8.5f} {time.time()-t0:8.5f}")
    avg_grad = 0.9 * avg_grad + 0.1 * grad

    w = w - learning_rates[n]*jnp.sign(avg_grad)

    t = time.time()-t0
    t = round(t * 1000, 5)

    if len(params)==1:
        params = unflatten(w)[0]
        # print(f" {t} | {params[0]} | {params[1]} | {params[2]} |
{avg_loss} ")
        return params
    else:
        params = unflatten(w)
        # print(f" {t} | {params[0]} | {params[1]} | {params[2]} |
{avg_loss} ")
        return unflatten(w)

def net(context,a,b,c):
    return a + b*context + c*context**2

def loss(pred, output):
    return (pred - output)**2

# arrays of batchsize, max_iter, learning_rates
batchsizes = [1,10,100,100,100]
max_iters = 1001
learning_rates = [[.1],[.1],[.1],[.01], [.1,.01]]

print(" Time (ms) | a | b | c |
final smooth loss")
print("-----|-----|-----|-----|-----
-----")

for batchsize, learning_rate in zip(batchsizes, learning_rates):
    key = jax.random.PRNGKey(0)
    a = jax.random.normal(key, 1)
    b = jax.random.normal(key, 1)
    c = jax.random.normal(key, 1)
    params = [a,b,c]

    gen = get_data_generator(batchsize)
    key = jax.random.PRNGKey(0)

    params = signed_gradient_descent(net, loss, gen, max_iters,
learning_rate, *params)
    print(f" {params} ")

```

Time (ms)	a	b	c	final
smooth loss				
-----	-----	-----	-----	-----

2605.96681	[3.6941566]	[7.0941534]	[-0.9058423]	
303.3845520019531				
2494.78602	[3.4941568]	[9.494158]	[-0.70584226]	
66.2212905883789				
2495.23616	[2.6941576]	[10.294161]	[-0.5058422]	
50.71369552612305				
2482.02205	[5.504189]	[8.924267]	[-0.59584194]	
112.87616729736328				
2494.96484	[3.0341573]	[10.014162]	[-0.6858425]	
31.885662078857422				

Question 3

1. It seems that smaller batch size (e.g. 1) results in slightly higher running time than bigger batch sizes (e.g. 100). Slower learning rates similarly result in slightly higher running time as well, although all are not significant changes.
2. It seems that larger batches and slower learning rates result in better final parameters, although not entirely consistent.
3. The final smooth loss seems to decrease with larger batchsizes. It also seems to work best with mixed learning rates, but in general, it seems that the final loss is lower with slower learning rate

Question 4

```

chars = ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9', 'a', 'b',
'c', 'd', 'e', 'f', 'g', 'h', 'i', 'j', 'k', 'l', 'm', 'n', 'o', 'p',
'q', 'r', 's', 't', 'u', 'v', 'w', 'x', 'y', 'z', 'A', 'B', 'C', 'D',
'E', 'F', 'G', 'H', 'I', 'J', 'K', 'L', 'M', 'N', 'O', 'P', 'Q', 'R',
'S', 'T', 'U', 'V', 'W', 'X', 'Y', 'Z', '!', '?', ':', '"', "'", '+',
',', '.', ' ', '\n']
start = [55, 17, 14, 70, 51, 27, 24, 19, 14, 12, 29, 70, 42, 30, 29,
14, 23,
11, 14, 27, 16, 70, 40, 37, 24, 24, 20, 70, 24, 15, 70, 55, 17,
14,
70, 40, 25, 18, 12, 70, 24, 15, 70, 54, 10, 30, 21, 68, 70, 11,
34,
70, 58, 18, 21, 21, 18, 10, 22, 70, 38, 21, 14, 10, 31, 14, 27,
70,
58, 18, 21, 20, 18, 23, 28, 24, 23]
res = ''.join([chars[c] for c in start])
print(res)

```

The Project Gutenberg EBook of The Epic of Saul, by William Cleaver Wilkinson

Question 5

```
import jax.numpy as jnp

data = jnp.load('data.npz', mmap_mode='r')['data']

def get_data_generator(context_size, batch_size):
    def getbatch(key):
        key, subkey = jax.random.split(key)
        start = jax.random.randint(subkey, shape=batch_size, minval=0,
maxval=len(data)-context_size)
        indices = start[:,None] + jnp.arange(context_size)[None,:]
        context = data[indices] # (batchsize x context_size)
        next = data[start+context_size] # (batchsize,)
        return context, next
    return getbatch

key = jax.random.PRNGKey(0)
getbatch = get_data_generator(50, 5)
context, next = getbatch(key)

print("Shape of context: ", context.shape, "Shape of next: ",
next.shape)

print("context: ", context)

print("next: ", next)

Shape of context: (5, 50) Shape of next: (5,)
context: [[24 23 28 25 18 27 10 12 34 68 71 14 33 29 14 23 13 14 13
70 11 34 70 29
17 14 70 36 12 29 28 70 24 15 70 1 8 8 7 68 70 1 8 9 1 68 70
10
23 13]
[18 29 14 13 70 54 29 10 29 14 28 70 12 24 25 34 27 18 16 17 29 70 18
23
70 29 17 14 28 14 70 32 24 27 20 28 68 70 28 24 70 29 17 14 70 41 24
30
23 13]
[28 69 70 44 29 70 18 28 70 25 27 24 11 10 11 21 14 70 29 17 10 29 70
39
30 27 17 10 22 71 32 10 28 70 27 14 10 21 21 34 70 22 24 27 14 70 18
23
15 21]
[23 70 29 17 14 70 56 23 18 29 14 13 70 54 29 10 29 14 28 68 70 32 14
70
13 24 70 23 24 29 70 12 21 10 18 22 70 10 70 27 18 16 17 29 70 29 24
```

```

70
 25 27]
[10 28 70 32 14 21 21 69 65 71 71 65 44 70 32 10 28 70 16 24 18 23 16
70
 29 24 70 10 28 20 70 34 24 30 70 29 24 70 21 14 29 70 22 14 70 29 10
20
 14 70]]
next: [70 10 30 14 10]

```

a) context shape: (5, 50). (batchsize, context_size) next shape: (5,) (batchsize)

b) context contains the batched numerical data that act as the input ("previous knowledge") to predict the target values contained in next.

c) the key is split into 2 subkeys, one for the context and one for next, such that we don't reuse the same random variables when generating.

question 6

```

from jax.nn import logsumexp

def loss(pred, next):
    num_characters = len(chars)
    assert pred.shape == (num_characters,)
    assert next.shape == ()
    log_probs = pred - logsumexp(pred)
    assert log_probs.shape == (num_characters,)
    next_onehot = jax.nn.one_hot(next, num_classes=num_characters)
    assert next_onehot.shape == (num_characters,)
    out = -jnp.sum(log_probs * next_onehot)
    assert out.shape == ()
    return out

def constant_net(context, b):
    # input context is a 1-D array of size context_size
    # each entry is an index between 0 and num_characters
    # these represent the most recent characters
    (context_size,) = context.shape
    (num_characters,) = b.shape

    # [do stuff.]

    # predict a constant vector
    pred = b

    assert pred.shape == (num_characters,)
    return pred

```

```

batchsize = 4096
context_size = 32
iters = 10000
learning_rates = [.001, .0001]
b = jnp.zeros(len(chars))

params = [b]

gen = get_data_generator(context_size, batchsize)

params = signed_gradient_descent(constant_net, loss, gen, iters,
learning_rates, *params)

```

iter	l.r.	loss	(smooth)	(avg)	time
0	0.00100	4.27667	4.27667	4.27667	0.02772
500	0.00100	3.75796	3.84823	4.00168	2.72620
1000	0.00100	3.42981	3.48416	3.79038	5.43914
1500	0.00100	3.28311	3.28676	3.63863	8.39022
2000	0.00100	3.18296	3.18663	3.53152	11.02354
2500	0.00100	3.13428	3.14404	3.45595	13.64159
3000	0.00100	3.13118	3.13523	3.40261	16.30306
3500	0.00100	3.14047	3.13310	3.36422	19.26403
4000	0.00100	3.10767	3.13320	3.33532	21.99465
4500	0.00100	3.11254	3.13212	3.31281	24.82710
5000	0.00010	3.12061	3.13259	3.29482	27.55235
5500	0.00010	3.11545	3.13263	3.28018	30.22116
6000	0.00010	3.14931	3.13334	3.26787	33.03475
6500	0.00010	3.15036	3.13409	3.25748	35.96461
7000	0.00010	3.14825	3.13339	3.24855	38.84472
7500	0.00010	3.11208	3.13300	3.24085	41.47183
8000	0.00010	3.14472	3.13200	3.23414	44.18035
8500	0.00010	3.11545	3.13456	3.22823	47.28741
9000	0.00010	3.14127	3.13248	3.22299	50.03408
9500	0.00010	3.13086	3.13249	3.21829	52.93438

Question 7

```

char2int = dict(zip(chars, range(len(chars))))

def generate_char(net, context, key, *params):
    pred = net(context, *params);
    assert pred.shape == (num_characters,)
    out = jax.random.categorical(key, pred);
    assert out.shape == ()
    return out

def generate(net, context_str, context_size, num_char, *params):
    context = [char2int[c] for c in context_str]

```

```

key = jax.random.PRNGKey(1)
for i in range(num_char):
    key, subkey = jax.random.split(key)
    my_context = jnp.array(context[-context_size:])
    c = generate_char(net, my_context, subkey, *params)
    context.append(int(c))

out = ''.join([chars[i] for i in context])

return out

b_constant = params
start = "STUDENT:\nI have searched the skies and found...\n\nTEACHER:\nYes? What have you found?\n\nSTUDENT\nI have found a fact, a fact most excellent.\n"
num_characters = len(chars)
print(generate(constant_net, start, context_size, 500, b_constant))

STUDENT:
I have searched the skies and found...

TEACHER:
Yes? What have you found?

STUDENT
I have found a fact, a fact most excellent.
c rtttu
e e
lv ib,grunek el iertah d
pauo goghmf e lrC"ac ty
eb tte biic,nitn io hp
oooiotttracvenn ern..ch I uen th ee
ha alelr woteInupo t r w
ro"e.lr
b ubDepegyttloi.hsinnceOfimtw d
wc,ernreluebn tnl dpcer sis ds reo "doh a ne c.satsfnm lsa rnf
rioefhaorcodees heotG e ao
oh tle dmhe eC ec sd2c,aride eei Th?ordr
Paehaeel l ud 0hrsoiad e eiavewi mr
dtf s fyahiSs eIn p oic olo
d r a Re mmth r leSess dto seia ielmrt hedluui2e s
liy ,frocn
fmce g dBndctud .iod ,agty l .L shEp

```

Question 8

```
import numpy as np
```



```

def linear_net(context, b, W):
    # W, b = params[0], params[1]
    # print('w: ', W)
    # print('b: ', b)

    (context_size,) = context.shape

    (context_size2, num_characters, num_characters2) = W.shape
    assert context_size == context_size2
    assert num_characters == num_characters2

    context_onehot = jax.nn.one_hot(context,
num_classes=num_characters)
    assert context_onehot.shape == (context_size, num_characters)

    # We need to multiply context_onehot with W along the context_size
dimension

    pred = jnp.einsum('ij,ijk->k', context_onehot, W) + b

    assert pred.shape == (num_characters,)
    return pred

batchsize = 4096
context_size = 32
iters = 10000
learning_rates = [.001, .0001]
num_characters = len(chars)

b_linear = jnp.zeros(num_characters)
W_linear = jnp.array(.01*np.random.randn(context_size, num_characters,
num_characters))

params = [b_linear, W_linear]

gen = get_data_generator(context_size, batchsize)

params = signed_gradient_descent(linear_net, loss, gen, iters,
learning_rates, *params)

```

iter	l.r.	loss	(smooth)	(avg)	time
0	0.00100	4.28394	4.28394	4.28394	0.16214
500	0.00100	2.56146	2.63390	2.84003	7.86799
1000	0.00100	2.37643	2.39678	2.64487	15.21716
1500	0.00100	2.33781	2.31454	2.54092	22.31885
2000	0.00100	2.29808	2.28205	2.47820	29.77769
2500	0.00100	2.26396	2.27183	2.43770	37.18627
3000	0.00100	2.28718	2.27311	2.41000	44.95227

3500	0.00100	2.24690	2.27569	2.39074	52.18431
4000	0.00100	2.29151	2.28119	2.37678	59.62810
4500	0.00100	2.28138	2.28403	2.36655	67.86923
5000	0.00010	2.28246	2.29063	2.35896	75.23881
5500	0.00010	2.30287	2.28773	2.35269	83.16335
6000	0.00010	2.29197	2.28485	2.34715	93.40257
6500	0.00010	2.35776	2.28828	2.34241	106.17691
7000	0.00010	2.31598	2.28488	2.33819	114.35264
7500	0.00010	2.24745	2.28746	2.33473	125.59524
8000	0.00010	2.26777	2.28614	2.33171	135.56565
8500	0.00010	2.27731	2.28815	2.32907	143.07084
9000	0.00010	2.32643	2.28784	2.32678	152.37469
9500	0.00010	2.31722	2.28497	2.32467	159.80256

Question 9

```
p = params
start = "STUDENT:\nI have searched the skies and found...\n\nTEACHER:\nYes? What have you found?\n\nSTUDENT\nI have found a fact, a fact most excellent.\n"
```

```
print(generate(linear_net,start,context_size,500, *p))
```

STUDENT:
I have searched the skies and found...

TEACHER:
Yes? What have you found?

STUDENT
I have found a fact, a fact most excellent.

of the s bllige the Amby whe probe d
"Prup of hmuthare thacllay
d were or canit the pestory: "Arcutwa cherragr at tenginee,
Ithale, as woted upolt haverobe.

N CLDecaly tiol. Sin centim, of caringelleby tho dpaknes somes aiof
dof aree cosstichm ls orne tion anorsedens he ther ablotele! meer ond
cow ton ine beem The rom
Painatel ufuthengsoind coniaq witho
d mante foreple ton proflyono
Mack are, mather theessadlougennin sizemethind uinne s inis ffrocher
fice of
Bricturn. He tasty he wall be

Question 10

```

def mlp_net(context, b, c, W, V):
    assert context.shape == (context_size,)
    (context_size,) = context.shape
    (num_characters,) = b.shape
    (num_hidden,) = c.shape
    (num_characters, num_hidden) = W.shape
    (context_size, num_characters, num_hidden) = V.shape

    context_onehot = jax.nn.one_hot(context,
num_classes=num_characters)
    assert context_onehot.shape == (context_size, num_characters)

    # [do stuff]

    # first layer
    h1 = jnp.einsum('ij,ijk->k', context_onehot, V) + c
    l1 = jax.nn.relu(h1)

    pred = b + W @ l1
    assert pred.shape == (num_characters,)
    return pred

batchsize = 4096
context_size = 32
iters = 10000
learning_rates = [.001, .0001]
num_characters = len(chars)

num_hidden = 500
b_mlp = jnp.zeros(num_characters)
c_mlp = jnp.zeros(num_hidden)
W_mlp = jnp.array(.01*np.random.randn(num_characters, num_hidden))
V_mlp = jnp.array(.01*np.random.randn(context_size, num_characters,
num_hidden))

params = [b_mlp, c_mlp, W_mlp, V_mlp]

gen = get_data_generator(context_size, batchsize)

params = signed_gradient_descent(mlp_net, loss, gen, iters,
learning_rates, *params)

```

iter	l.r.	loss (smooth)	(avg)	time
0	0.00100	4.27603	4.27603	0.79735
500	0.00100	1.90563	2.00908	34.99069
1000	0.00100	1.78632	1.80122	71.84578
1500	0.00100	1.74954	1.73723	112.49686
2000	0.00100	1.74241	1.70301	156.91322

2500	0.00100	1.62852	1.67725	1.85782	194.31828
3000	0.00100	1.66574	1.65953	1.82540	233.58647
3500	0.00100	1.65879	1.64913	1.80061	270.97481
4000	0.00100	1.63123	1.64054	1.78076	310.07613
4500	0.00100	1.58940	1.62880	1.76425	349.32883
5000	0.00010	1.62587	1.62300	1.75032	386.42197
5500	0.00010	1.64957	1.59733	1.73701	425.49937
6000	0.00010	1.61687	1.58785	1.72477	463.36210
6500	0.00010	1.62156	1.58363	1.71394	499.46647
7000	0.00010	1.58306	1.57478	1.70408	538.08612
7500	0.00010	1.52659	1.57148	1.69539	577.97836
8000	0.00010	1.56541	1.56741	1.68757	614.37243
8500	0.00010	1.56117	1.56740	1.68047	649.60792
9000	0.00010	1.54836	1.56470	1.67399	685.35523
9500	0.00010	1.62733	1.55928	1.66808	720.54154

Question 11

```
p = params
start = "STUDENT:\nI have searched the skies and found...\n\nTEACHER:\nYes? What have you found?\n\nSTUDENT\nI have found a fact, a fact most excellent.\n"
```

```
print(generate(mlp_net,start,context_size,500, *p))
```

STUDENT:

I have searched the skies and found...

TEACHER:

Yes? What have you found?

STUDENT

I have found a fact, a fact most excellent.

To rettire liver, restor elexpress delierso of the Daymarce! The difficulting, it will how to the cuttains one. Wat feeling mark, took the Vedaup on a swork at leagure which it grain confided two, the lutby that parts, succeeded by him he cousts of liverning obland convested the family were began anced boaring being to rell and and leud andsomal conicient of the Volvania, satingphanced of discard, mather leters who said, as immorded upon the like, by heart clught becauded. The stock fastshop

Question 12

```
def dbl_net(context, b, c, d, W, V, U):
    assert context.shape == (context_size,)
```

```

    context_onehot = jax.nn.one_hot(context,
num_classes=num_characters)
    assert context_onehot.shape == (context_size, num_characters)

    # [do stuff]
    h1 = jnp.einsum('ijk,ij->k', V, context_onehot) + c
    l1 = jax.nn.relu(h1)

    # second layer
    h2 = jnp.einsum('k,jk->k', l1, U) + d
    l2 = jax.nn.relu(h2)

    pred = b + W @ l2
    assert pred.shape == (num_characters,)

    return pred

batchsize = 4096
context_size = 32
iters = 10000
learning_rates = [.001, .0001]
num_characters = len(chars)

num_hidden = 500
b_dbl = jnp.zeros(num_characters)
c_dbl = jnp.zeros(num_hidden)
d_dbl = jnp.zeros(num_hidden)
W_dbl = jnp.zeros((num_characters, num_hidden))
V_dbl = jnp.array(.01*np.random.randn(context_size, num_characters,
num_hidden))
U_dbl = jnp.array(.01*np.random.randn(num_hidden, num_hidden))

params = [b_dbl, c_dbl, d_dbl, W_dbl, V_dbl, U_dbl]

gen = get_data_generator(context_size, batchsize)

params = signed_gradient_descent(dbl_net, loss, gen, iters,
learning_rates, *params)

```

iter	l.r.	loss	(smooth)	(avg)	time
0	0.00100	4.27667	4.27667	4.27667	0.39515
500	0.00100	1.91986	1.98242	2.21543	52.13944
1000	0.00100	1.84668	1.86097	2.04942	102.83609
1500	0.00100	1.85455	1.82285	1.97715	152.54611
2000	0.00100	1.82844	1.80457	1.93525	202.87642
2500	0.00100	1.76737	1.79316	1.90756	253.00644
3000	0.00100	1.77197	1.78332	1.88719	306.35678
3500	0.00100	1.78004	1.78175	1.87214	356.96288
4000	0.00100	1.76319	1.77826	1.86052	407.42856

4500	0.00100	1.74391	1.77386	1.85109	456.78786
5000	0.00010	1.76011	1.77274	1.84338	504.86405
5500	0.00010	1.76384	1.72043	1.83286	554.39541
6000	0.00010	1.74332	1.71038	1.82284	601.72159
6500	0.00010	1.74147	1.70693	1.81395	649.82745
7000	0.00010	1.72367	1.69977	1.80580	700.99020
7500	0.00010	1.64946	1.69591	1.79862	751.55216
8000	0.00010	1.67967	1.69293	1.79215	804.38640
8500	0.00010	1.70911	1.69330	1.78629	854.83338
9000	0.00010	1.68200	1.69000	1.78097	903.49232
9500	0.00010	1.74276	1.68659	1.77616	957.79651

Question 13

```
p = params
start = "STUDENT:\nI have searched the skies and found...\n\nTEACHER:\nYes? What have you found?\n\nSTUDENT\nI have found a fact, a fact most excellent.\n"
```

```
print(generate(dbl_net,start,context_size,500, *p))
```

STUDENT:

I have searched the skies and found...

TEACHER:

Yes? What have you found?

STUDENT

I have found a fact, a fact most excellent.

To rettle s breation oundral, for the deticalong gome that Came bar
But Jew Timc, it with and to the cuncuess on. You duen the mark,
though world upon a way
routly bevered by thod."

"He word,work at the Cebjects Rock, with discendiders
where.

Scrusist ruserious arroodeds here was overwanted a both cown't wite
been breard'

One in where so, I canifus into
dufter Royaging to pastic plopt they be maintailed scant, such found
other to him stilike, by he had lightened up of detage man which
p

Question 14

```

batchsize = 4096
context_size = 32
iters = 10000
learning_rates = [.001, .0001]
num_characters = len(chars)

num_hidden = 2500
b_dbl = jnp.zeros(num_characters)
c_dbl = jnp.zeros(num_hidden)
d_dbl = jnp.zeros(num_hidden)
W_dbl = jnp.zeros((num_characters, num_hidden))
V_dbl = jnp.array(.01*np.random.randn(context_size, num_characters,
num_hidden))
U_dbl = jnp.array(.01*np.random.randn(num_hidden, num_hidden))

params = [b_dbl, c_dbl, d_dbl, W_dbl, V_dbl, U_dbl]

gen = get_data_generator(context_size, batchsize)

params = signed_gradient_descent(dbl_net, loss, gen, iters,
learning_rates, *params)

```

iter	l.r.	loss (smooth)	(avg)	time
0	0.00100	4.27667	4.27667	0.53196
500	0.00100	1.95496	2.00869	231.11978
1000	0.00100	1.90694	1.91072	2.06173
1500	0.00100	1.92144	1.90963	2.01003
2000	0.00100	1.89641	1.87344	1.97817
2500	0.00100	1.83292	1.86724	1.95542
3000	0.00100	1.85139	1.85567	1.93868
3500	0.00100	1.86184	1.85394	1.92677
4000	0.00100	1.85162	1.86234	1.91908
4500	0.00100	1.80985	1.87569	1.91336
5000	0.00010	1.84876	1.86237	1.90806
5500	0.00010	1.82357	1.78007	1.89734
6000	0.00010	1.79365	1.76830	1.88672
6500	0.00010	1.79828	1.76144	1.87713
7000	0.00010	1.76979	1.75612	1.86842
7500	0.00010	1.71450	1.75216	1.86071
8000	0.00010	1.72298	1.74601	1.85365
8500	0.00010	1.76215	1.74602	1.84727
9000	0.00010	1.74628	1.74563	1.84152
9500	0.00010	1.78409	1.73857	1.83623

Question 15

p = params

```
start = "STUDENT:\nI have searched the skies and found...\n\nTEACHER:\nYes? What have you found?\n\nSTUDENT\nI have found a fact, a fact most excellent.\n"
```

```
print(generate(dbl_net,start,context_size,500, *p))
```

STUDENT:
I have searched the skies and found...

TEACHER:
Yes? What have you found?

STUDENT
I have found a fact, a fact most excellent.

The they we bwlitied to the slexervicil
No goghmen. We Ccame by the Jew bish, it his passon to the know
serval hak benother,
have elrast other polit way robe. Befoue
boways ind."335 Oblowers," carred lief "And Rpcers, sumpassed the
barness.

Scrust!" And the fairrs Mrss herood for beale on the Caprod William
He may Mord'
Pheianes leuded thou will if still sonfented angios to pastic of the
Charless the readess who sevent in mull had in ploy
like, bove o't clugidenchaud mind togeth was, shop